

## **Unit 1: AI Agents (35%)**

### **Exam Practice Questions**

**1. Which statement accurately describes the fundamental difference between context variables and conversation variables?**

- A) Context variables (prefixed \$Context) are read-only after session initialization; conversation variables are created via action outputs in Agent Builder to store data from actions during the conversation.
- B) Context variables are created during actions and can be modified; conversation variables are set at the start of a session and remain static.
- C) Context variables store conversation history updated by the LLM; conversation variables are system metadata that is not accessible during chats.

**2. An Agentforce Specialist needs to verify a user and then use the verified customer ID in subsequent actions within Agentforce. How can the specialist efficiently pass the verified customer ID between actions?**

- A) Use a context variable.
- B) Use a custom conversation variable.
- C) Use a filter variable.

**3. An Agentforce Specialist wants to ensure that certain actions in Agentforce only execute if a customer is verified. What should the specialist use to enforce this condition?**

- A) Context variable.
- B) Custom conversation variable.
- C) Filter.

**4. A Service Agent is now unable to help customers with issues related to a new product replacement process managed by a custom Product Replacement object. Which Agentforce Agent User change must be implemented?**

- A) The permission set assigned to the Agent User needs Read access to the custom Product Replacement object.
- B) The profile assigned to the Agentforce Agent User needs AI training permission.
- C) The permission set group assigned to the Agent User needs to grant access to the Product Replacement flow.

### **Secondary Quiz Assessments**

- **Discover Agentforce:**

1. What is a topic classification? (Correct: B - A group of user utterances based on predefined descriptions).
2. Which of these can serve as guardrails for agents? (Correct: E - Natural-language instructions and built-in security features).

• **Customize Your Agents:**

1. Which type of customization is best for complex automation, like retrieving data from a third-party service as part of a prompt? (Correct: C - Apex).
2. What do flows allow AI agents to do? (Correct: A - Make changes to your Salesforce data).

• **Guardrails and Trust:**

1. How can you adjust the agent's behavior in Agentforce? (Correct: C - Adjust topic instructions and actions).
2. True or false: Topic classification is used to map user inputs to specific topics. (Correct: A - True).

• **Responsible and Ethical Agents:**

1. How does Salesforce approach red-teaming for its AI products? (Correct: B - By using both manual and automated red-teaming methods).
2. True or false: Continuous improvement and innovation are crucial for maintaining the integrity of AI systems. (Correct: A - True).

• **Explore Agentforce:**

1. What's the primary difference between standard and custom actions? (Correct: A - Standard are predefined; custom are built for specific business needs).
2. Which describes the role of topics? (Correct: B - Provide agent with a set of actions and instructions).

• **Meet the Reasoning Engine:**

1. What is the primary role of the reasoning engine? (Correct: B - To interpret user requests and classify them into topics).
  2. What does it do when it finds and launches the right topics and actions? (Correct: C - It retrieves and provides detailed, relevant information).
- 

## **Unit 2: Prompt Engineering (20%)**

### **Exam Practice Questions**

**1. A specialist needs a prompt template that generates a summary of only open cases for an account. The grounding must return only open cases directly. Which approach best meets this?**

- A) Ground with Apex Merge Fields.
- B) Ground with Flow Merge Fields.
- C) Ground with Record Merge Fields.

**2. What is a key benefit of grounding prompt templates with CRM data in Prompt Builder?**

- A) It scores content on toxicity.
- B) It automatically updates Salesforce records.
- C) It provides the LLM with context to create personalized responses.

**3. How should a marketing team ensure a Flex prompt template generates personalized event invitations for each contact?**

- A) Include instructions in the prompt template.
  - B) Manually input contact details.
  - C) Integrate the template with dynamic CRM data fields for automatic customization.
- 

**Unit 3: Data Cloud for Agentforce (20%)**

**Exam Practice Questions**

**1. What is the primary purpose of using retrievers in Data 360?**

- A) To store unstructured data.
- B) To create and manage search indexes.
- C) To return relevant data from the vector database to augment a prompt.

**2. Which action should a company take to enhance LLM accuracy by incorporating relevant knowledge from various sources using retrievers?**

- A) Add active retrievers to prompt templates to retrieve knowledge.
- B) Use related lists and Apex to dynamically generate search text.
- C) Deactivate default retrievers and rely on custom retrievers.

**Secondary Quiz Assessments**

• **Agentforce Library Basics:**

1. What is the main challenge in integrating unstructured data into AI features? (Correct: B - It is harder to search but contains valuable insights).
2. What occurs at runtime when a user asks a question in an Agentforce Data Library? (Correct: B - User's query is added to prompt template, and retriever fetches relevant data).

• **Build a File-Based Data Library:**

1. What should you do to test out changes made to a data library? (Correct: B - Restart the preview conversation and chat with your agent).
2. Where can you find which sources of information your agent uses for its responses? (Correct: B - In the agent reasoning window).

• **Data Cloud Overview Unit:**

1. The section makes up what percentage of the exam? (Correct: C - 20%).
  2. Which key topic does this section cover? (Correct: B - Improving agent response with unstructured data using chunking and indexing).
- 

**Unit 4: Development Lifecycle (20%)**

**Exam Practice Questions**

**1. Universal Containers needs to automate regression testing within CI pipelines. Which method should the team use?**

- A) Use Agentforce DX or Testing API to run scripted tests via CLI or REST API.
- B) Manually rerun test cases daily.
- C) Defer regression testing until an error is reported.

**2. Universal Containers wants to fast-track its testing process and use AI-generated test cases. Which steps should they follow?**

- A) Deploy to production without testing.
- B) Use Testing Center, add context, generate AI test cases, and upload a CSV.
- C) Manually rerun test cases daily without AI.

**3. What is a critical predeployment step to ensure a successful Agentforce deployment from sandbox to production?**

- A) Ensure Target Org has same profiles.
- B) Verify Target Org has same data records.
- C) Verify Target Org has valid licenses for Agentforce and ensure Einstein and Bots are enabled.

**4. Which step is essential for managing and monitoring Agentforce agent adoption using session tracing?**

- A) Enable session tracing and review logs in console.
- B) Enable session tracing and use Data 360 to build reports and dashboards.
- C) Enable session tracing and use built-in real-time metrics.

### **Secondary Quiz Assessments**

- **View Test Results:**

1. Which part of the test results does Actual Actions represent? (Correct: C - The actions the agent ran).
  2. Agentforce Builder allows for conducting which type of testing? (Correct: A - Manual testing).
- 

### **Unit 5: Multi-Agent Interoperability (5%)**

#### **Quiz Assessment**

**1. The Multi-Agent Interoperability section makes up what percentage of the certification exam?**

- A) 5%.

**2. Which key topic does the Multi-Agent Interoperability section cover?**

- B) Explain the purpose of Model Context Protocol (MCP) and its use cases.
-

## General Knowledge / Secondary Topics

### Autonomous Agents Quiz

- **Get to Know Autonomous Agents:** 1. Primary function? (B - Save time/resources, enhance experience). 2. Distinguish from chatbots? (D - Execute tasks independently).
- **Establish Trust:** 1. Guardrail to consider? (B - Provide human oversight for complex interactions). 2. Transparent about? (A - How/when customers interact with agent).
- **Transform Business:** 1. Benefit in customer support? (B - Human reps focus on complex issues). 2. Ensures data protection/toxicity detection? (B - Einstein Trust Layer).

### LLM & Salesforce Basics Quiz

- **LLM Basics:** 1. Variables learned during training? (D - Parameters). 2. Benefit of increasing parameters? (B - Improve human-like text generation).
- **Salesforce & LLMs:** 1. Need more than one LLM? (A - True). 2. Benefit of geo-aware model? (C - Greater control over data residency).

### 1. AI Agents (35%) - Missing Details

- **Permissions for Management:** To define users as **Service agent managers** so they can create and edit agents, you must assign them the **Manage Agentforce Service Agents** permission set.
- **The Difference Between Bots and Agents:** Unlike Einstein Bots, which are based on **complex, strictly defined conversational rules**, Agentforce is flexible and requires less configuration because it is **powered by an LLM**.
- **Topic Classification:** This is the specific trust pattern used to **reduce hallucinations** by mapping user inputs to specific, governed topics.
- **Ethical Guardrails:** While the Einstein Trust Layer provides security, **topic instructions** are the primary way to define **ethical boundaries** (e.g., "Always check if the customer is verified...").

### 2. Prompt Engineering (20%) - Missing Numerical Limits

The sources list specific numerical limits that are critical for a Specialist to know when building templates:

- **Maximum template size:** 128,000 characters.
- **Maximum number of template versions:** 50.
- **Maximum number of merge fields:** 50.
- **Maximum number of Flow/Apex/Related List merge fields:** 5 per template.
- **File Upload Limits:** 10 images or 100 PDF pages (total 15MB) per request.

### 3. Data Cloud for Agentforce (20%) - Process Details

- **Unstructured Data Types:** Approximately **90% of enterprise data** is unstructured (videos, emails, audio, PDFs), which is harder to search but contains vital customer sentiment.

- **Setup Automation:** Creating a data library automatically pushes **data streams to Data 360**, maps data objects, and creates a **search index and retriever**.
- **Runtime Logic:** At runtime, the user's query is added to the **prompt template**, which then calls the **retriever** to fetch data from the **search index**.

#### 4. Development Lifecycle (20%) - Testing Limits & Tools

- **Testing Center Limits:** You can run up to **10 test jobs at once** in a 10-hour timeframe, with up to **1,000 test cases per test**.
- **CLI Support:** The Salesforce CLI now supports **Data Cloud and Agentforce** metadata for deployments.
- **Testing Ethics:** "Red-teaming" is the structured process of probing AI for **harmful capabilities** or **prompt injections**.
- **Test-Driven Development (TDD):** The sources suggest you can create test criteria *first*, run the tests, and then build agent functionality to get those tests to pass.

#### 5. Multi-Agent Interoperability (5%) - Key Use Cases

- **Agent API for Downstream Sessions:** The Agent API is specifically useful for **starting a downstream agent's session** and passing a verified customer ID as a read-only context variable [Multi-Agent Interoperability Section].
- **Interoperability Benefit:** These protocols allow agents to **share context** and work together across different vendors.

---

**Mentor's Final Tip:** The exam will often ask "What is the best way to...?"

- If the goal is **speed and exact match**, the answer is **Keyword Search**.
- If the goal is **handling misspellings/similarity**, the answer is **Vector/Semantic Search**.
- If the goal is **maximum reliability**, the answer is **Hybrid Search**.

You have now reviewed all the material from the sources provided. To solidify this, imagine the **Einstein Trust Layer** as a **Security Detail** for a celebrity (your data). They check the "guest list" (permissions), "mask" the celebrity's face in public (data masking), and "zero-out" any memories the paparazzi (third-party LLMs) might try to keep.